

Learning from Data  
4. Classification: Nearest Neighbor

© Prakash Ishwar  
Spring 2017

# Classification

- Supervised (predictive) learning: given **examples** with **labels**, predict labels for all **unseen** examples
  - Classification:
    - label = category,
    - $\mathbf{y} \in \mathcal{Y} = \{1, \dots, m\}$ ,  $m$  = number of classes
    - $\ell(\mathbf{x}, y, h) = 1(h(\mathbf{x}) \neq y)$ , Risk =  $P(Y \neq h(\mathbf{X})) = P(\text{Error})$




$\mathbf{x}$  = facial geometry features  
 $y$  = gender label

# k-Nearest Neighbor (NN) Classifier

- **Discriminative classifier:** only  $p(y|\mathbf{x})$  estimated
- Non-parametric: no parametric model for  $p(y|\mathbf{x})$
- Example of memory-based learning
- Need 2 ingredients to specify classifier:
  1.  $k$  : number of nearest neighbors, typically chosen to be not a multiple of  $m$  (the number of classes)
  2.  $\text{dist}: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_+$ , a measure of “nearness”  $\text{dist}(\mathbf{x}, \mathbf{x}')$  between  $\mathbf{x}$  and  $\mathbf{x}'$ 
    - typically,  $\text{dist}$  is chosen to be a metric, i.e., it is positive, definite, i.e.,  $= 0$  if, and only if,  $\mathbf{x} = \mathbf{x}'$ , symmetric, i.e.,  $\text{dist}(\mathbf{x}, \mathbf{x}') = \text{dist}(\mathbf{x}', \mathbf{x})$ , and satisfies the triangle inequality, i.e.,  $\text{dist}(\mathbf{x}_1, \mathbf{x}_3) \leq \text{dist}(\mathbf{x}_1, \mathbf{x}_2) + \text{dist}(\mathbf{x}_2, \mathbf{x}_3)$  for any three points

# k-Nearest Neighbor Classifier

- Classifier description in words:  $h_{k\text{-NN}}(\mathbf{x})$  = most abundant label (majority vote) among the labels of the  $k$  nearest training examples of  $\mathbf{x}$  (breaking ties in some way)



# k-Nearest Neighbor Classifier

- **Formal description:** Given labeled training data  $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$  and a test point  $\mathbf{x}$
- Let  $\{(\mathbf{x}_{(1)}, y_{(1)}), \dots, (\mathbf{x}_{(n)}, y_{(n)})\}$  be a re-ordering of training data such that

$$d(\mathbf{x}, \mathbf{x}_{(1)}) \leq d(\mathbf{x}, \mathbf{x}_{(2)}), \dots, d(\mathbf{x}, \mathbf{x}_{(n)})$$

- Then,

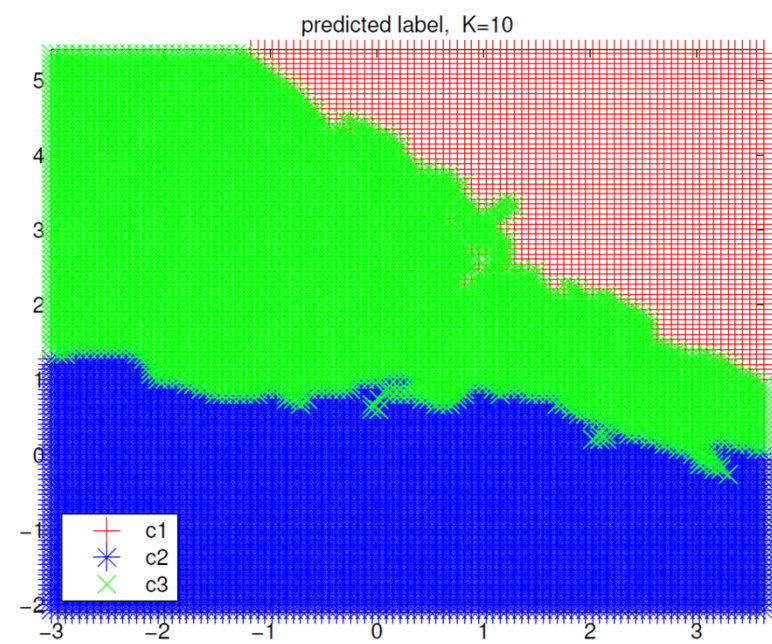

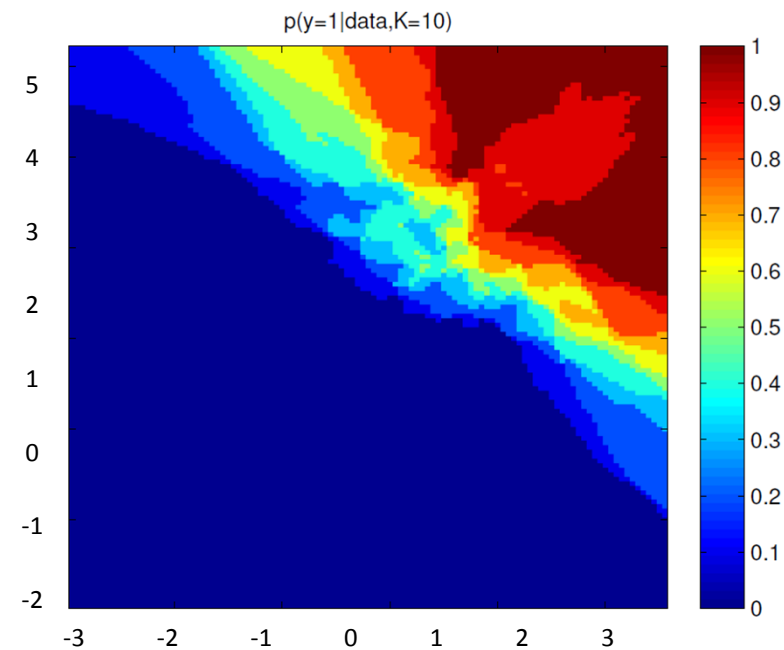

$$h_{k-\text{NN}}(\mathbf{x}) = \arg \max_{y=1, \dots, m} \underbrace{\sum_{j=1}^k 1(y_{(j)} = y)}_{\text{number of } k \text{ NNs of } \mathbf{x} \text{ with label } = y}$$

# k-Nearest Neighbor Classifier

- Discriminative model based interpretation:

$$\hat{p}_{k\text{-NN}}(y|\mathbf{x}) = \underbrace{\frac{1}{k} \sum_{j=1}^k 1(y_{(j)} = y)}_{\text{fraction of } k \text{ NNs of } \mathbf{x} \text{ with label } = y}$$

- This is a non-parametric estimate of  $p(y|\mathbf{x})$ , since the number of parameters =  $n$  grows with training data
- k-NN classifier = MPE/MAP rule with the above non-parametric estimate of  $p(y|\mathbf{x})$ .



# Common distance functions


- $\ell_p$  distance:

$$\|\mathbf{x} - \mathbf{x}'\|_p := \left( \sum_{i=1}^d |x_i - x'_i|^p \right)^{\frac{1}{p}}$$

- can prove that this is a norm-distance for all  $p \geq 1$
- Euclidean or  $\ell_2$  distance:  $\ell_p$  distance with  $p = 2$
- taxi-cab or city-block or Manhattan or  $\ell_1$  distance:  $\ell_p$  distance with  $p = 1$
- max norm or  $\ell_\infty$  distance: limit of  $\ell_p$  as  $p \rightarrow \infty$ . Can be shown to be equal to:  $\max_{1 \leq i \leq d} |x_i - x'_i|$
- Mahalanobis distance:  $\sqrt{(\mathbf{x} - \mathbf{x}')^T \Sigma^{-1} (\mathbf{x} - \mathbf{x}')}$  with  $\Sigma$  a positive definite square matrix



# Common distance functions



Contours of constant  $p$ -norms  
(distance from origin)  
in 2 dimensions ( $d=2$ ) for different  $p$

# Voronoi regions for 1-NN classifiers


- **Voronoi region** of a training point  $\mathbf{x}_j$  = all points in feature space that are closer to  $\mathbf{x}_j$  than to any other training point:

$$\mathcal{V}(\mathbf{x}_j) := \{\mathbf{x} : d(\mathbf{x}, \mathbf{x}_j) \leq d(\mathbf{x}, \mathbf{x}_{j'}), \forall j' \neq j\}$$

- $\Rightarrow$  The 1-NN classifier will assign all points in  $\mathcal{V}(\mathbf{x}_j)$  the same class label as that of  $\mathbf{x}_j$
- The **Voronoi tessellation** is a **partition** of the feature space into the Voronoi regions of the training points
- If the distance function is a norm induced by an inner product, e.g.,  $\|\mathbf{x}\| = \sqrt{\langle \mathbf{x}, \mathbf{x} \rangle}$ , then the Voronoi regions are **convex**, i.e., for any two points in the region, the entire line segment joining them is also entirely in the region.

# Voronoi regions for 1-NN classifiers

- Voronoi tessellation of the 2-dim plane induced by the 1-NN classifier based on **Euclidean distance**




# Voronoi regions for 1-NN classifiers

- Voronoi tessellation based on:

# Euclidean distance (convex)


versus

# Manhattan distance (non-convex)



# Voronoi regions for 1-NN classifiers

- 2-point Voronoi diagram for **Manhattan distance**



# Remarks

- Asymptotic performance guarantees:
  - as  $n \rightarrow \infty$ , risk of 1-NN classifier (for 0-1 loss)  $\leq 2R_{\text{Bayes}}$  for all “nice” data distributions and distance functions!
  - If as  $n \rightarrow \infty$ ,
    - $k_n \rightarrow \infty$  ... (ensures more training examples in majority vote)
    - $\frac{k_n}{n} \rightarrow 0$  ... (ensures the  $k_n$  NNs get closer to any test point)then, risk of  $k_n$ -NN classifier (for 0-1) loss  $\rightarrow R_{\text{Bayes}}$  for all “nice” data distributions and distance functions!!
- In practice, the best value of  $k$  is determined via cross-validation

## Risk of 1-NN classifier (for 0-1 loss) $\leq 2R_{\text{Bayes}}$

**Result:** as  $n \rightarrow \infty$ ,

$$P(Y_{\text{test}} \neq h_{1-NN}(\mathbf{X}_{\text{test}})) \leq 2P(Y_{\text{test}} \neq h_{\text{Bayes}}(\mathbf{X}_{\text{test}}))$$

for all “nice” joint distributions  $p(y|\mathbf{x})p(\mathbf{x})$  and distance functions  $\text{dist}(\cdot, \cdot)$

**Intuition** for  $m = 2$  classes. Let  $\mathcal{Y} = \{0,1\}$ .

- Training set:  $(\mathbf{X}_1, Y_1), \dots, (\mathbf{X}_n, Y_n) \sim \text{IID } p(y|\mathbf{x})p(\mathbf{x})$
- Test point:  $(\mathbf{X}_{\text{test}}, Y_{\text{test}}) \sim p(y|\mathbf{x})p(\mathbf{x})$  independent of training set
- Let  $NN(\mathbf{x})$  = training feature closest to  $\mathbf{x}$  and let  $Y_{NN(\mathbf{x})}$  denote its label

## Risk of 1-NN classifier (for 0-1 loss) $\leq 2R_{\text{Bayes}}$

- **Key observation 1:** if  $\mathbf{x}$  is within the interior of the support of  $p(\mathbf{x})$  and  $p(\mathbf{x})$ ,  $\text{dist}(\cdot, \cdot)$  are “nice”, then as  $n \rightarrow \infty$ ,  $NN(\mathbf{x}) \rightarrow \mathbf{x}$  almost surely:

$$\forall \epsilon > 0, P(\text{dist}(\mathbf{x}, NN(\mathbf{x})) \geq \epsilon) = (1 - \underbrace{P(\text{dist}(\mathbf{x}, \mathbf{X}) \geq \epsilon)}_{>0 \text{ since } \mathbf{x} \in \text{int supp}(p(\mathbf{x})) \text{ and } p(\mathbf{x}), \text{dist}(\cdot, \cdot) \text{ are nice}})^n \rightarrow 0$$

- **Key observation 2:** if  $p(y|\mathbf{x})$  is “nice” then whenever  $\mathbf{x}'$  is close to  $\mathbf{x}$  (as measured by  $\text{dist}(\cdot, \cdot)$ ) then  $p(y|\mathbf{x}')$  will be close to  $p(y|\mathbf{x})$

- **Key observation 3:**

$$P(Y_{\text{test}} \neq h_{\text{Bayes}}(\mathbf{X}_{\text{test}}) | \mathbf{X}_{\text{test}} = \mathbf{x}) = \min\{p(0|\mathbf{x}), p(1|\mathbf{x})\}$$

$$\text{since } h_{\text{Bayes}}(\mathbf{x}) = \arg \max_{y=0,1} p(y|\mathbf{x})$$



# Risk of 1-NN classifier (for 0-1 loss) $\leq 2R_{\text{Bayes}}$

- Proof-sketch:

$$\begin{aligned} & P(Y_{\text{test}} \neq h_{1\text{-NN}}(\mathbf{X}_{\text{test}}) | \mathbf{X}_{\text{test}} = \mathbf{x}) \\ &= P(Y_{\text{test}} \neq Y_{NN(\mathbf{x})} | \mathbf{X}_{\text{test}} = \mathbf{x}) \\ &= P(Y_{\text{test}} = 1 | \mathbf{X}_{\text{test}} = \mathbf{x}) P(Y_{NN(\mathbf{x})} = 0 | \mathbf{X}_{\text{test}} = \mathbf{x}) \quad \text{training and test samples} \\ &\quad + P(Y_{\text{test}} = 0 | \mathbf{X}_{\text{test}} = \mathbf{x}) P(Y_{NN(\mathbf{x})} = 1 | \mathbf{X}_{\text{test}} = \mathbf{x}) \quad \text{are independent} \\ &= p(1|\mathbf{x}) P(Y_{NN(\mathbf{x})} = 0 | \mathbf{X}_{\text{test}} = \mathbf{x}) + p(0|\mathbf{x}) P(Y_{NN(\mathbf{x})} = 1 | \mathbf{X}_{\text{test}} = \mathbf{x}) \\ &\approx p(1|\mathbf{x}) P(Y_{\text{test}} = 0 | \mathbf{X}_{\text{test}} = \mathbf{x}) + p(0|\mathbf{x}) P(Y_{\text{test}} = 1 | \mathbf{X}_{\text{test}} = \mathbf{x}) \quad \text{Key} \\ &= p(1|\mathbf{x}) p(0|\mathbf{x}) + p(0|\mathbf{x}) p(1|\mathbf{x}) \quad \text{observations} \\ &= 2p(0|\mathbf{x}) p(1|\mathbf{x}) \quad \text{1 and 2} \\ &\leq 2 \min\{p(0|\mathbf{x}), p(1|\mathbf{x})\} \quad \text{If } a, b \in [0,1] \Rightarrow ab \leq a \text{ and } ab \leq b \Rightarrow ab \leq \min\{a, b\} \\ &= 2P(Y_{\text{test}} \neq h_{\text{Bayes}}(\mathbf{X}_{\text{test}}) | \mathbf{X}_{\text{test}} = \mathbf{x}) \quad \text{Key observation 3} \end{aligned}$$

- The result follows by taking expectation over  $\mathbf{X}_{\text{test}}$

# Remarks

- **Useful observation:** Euclidean distances can be computed via inner products:  $\|\mathbf{x} - \mathbf{x}'\|_2^2 = \langle \mathbf{x}, \mathbf{x} \rangle + \langle \mathbf{x}', \mathbf{x}' \rangle - 2\langle \mathbf{x}, \mathbf{x}' \rangle$
- Let  $\mathbb{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]$ ,  $\text{diag}(A)$  = column vector of the main diagonal elements of square matrix  $A$ , and  $\mathbf{1}$  = column vector of all ones. Then the square Euclidean distance matrix (EDM) is given by:

$$\text{EDM}(\mathbb{X}) = \mathbf{1} * \text{diag}(\mathbb{X}^T \mathbb{X})^T - 2\mathbb{X}^T \mathbb{X} + \text{diag}(\mathbb{X}^T \mathbb{X}) * \mathbf{1}^T$$

$$\text{EDM}(\mathbb{X})(i, j) = \|\mathbf{x}_i - \mathbf{x}_j\|_2^2$$

- **Note:**  $\text{rank}(\text{EDM}(\mathbb{X})) \leq d + 2$

# Weighted nearest neighbor rules

- Observation: 
$$\begin{aligned} h_{k-NN}(\mathbf{x}) &= \arg \max_y \hat{p}_{k-NN}(y|\mathbf{x}) \\ &= \arg \max_y \sum_{j=1}^k \frac{1}{k} 1(y_{(j)} = y) \\ &= \arg \max_y \sum_{j=1}^n \underbrace{w(\mathbf{x}, \mathbf{x}_j)}_{\text{weights}} 1(y_j = y) \end{aligned}$$

where

$$w(\mathbf{x}, \mathbf{x}_j) = \begin{cases} \frac{1}{k} & \text{if } \mathbf{x}_j \in \{k\text{NN's of } \mathbf{x}\} \\ 0 & \text{otherwise} \end{cases}$$

- **Idea:** use alternative weight function  $w(\mathbf{x}, \mathbf{x}')$ : **nonnegative** and **decreases with increasing distance**  $\text{dist}(\mathbf{x}, \mathbf{x}')$ . Can be interpreted as a “**similarity**” **score**. Typically,

$$w(\mathbf{x}, \mathbf{x}') \propto e^{-\text{dist}(\mathbf{x}, \mathbf{x}')}$$

# Weighted nearest neighbor rules

Example:

$$w(\mathbf{x}, \mathbf{x}') = \frac{1}{\alpha\sqrt{2\pi}} e^{-\frac{1}{2\alpha^2}\|\mathbf{x}-\mathbf{x}'\|^2} = \frac{1}{\alpha} \psi\left(\frac{\|\mathbf{x}-\mathbf{x}'\|}{\alpha}\right), \psi(t) = \mathcal{N}(0,1)(t)$$


Here,  $\alpha$  is a **bandwidth parameter** which controls the **neighborhood size** like  $k$  in  $k$ -NN (small  $\alpha \Rightarrow$  small neighborhood)

- $$\hat{p}(\mathbf{x}|y) = \frac{1}{n_y} \sum_{j=1}^n \frac{1}{\alpha\sqrt{2\pi}} e^{-\frac{1}{2\alpha^2}\|\mathbf{x}-\mathbf{x}_j\|^2} 1(y_j = y), n_y = |\{j : y_j = y\}|$$
- is a **non-parametric estimate** of the conditional density  $p(\mathbf{x}|y)$
- called kernel-density estimate with the Gaussian “kernel” of bandwidth  $\alpha$
- Thus,
$$\hat{p}(y|\mathbf{x}) \propto \hat{p}(\mathbf{x}|y) \underbrace{\hat{p}(y)}_{\frac{n_y}{n}} = \frac{1}{n} \sum_{j=1}^n \frac{1}{\alpha\sqrt{2\pi}} e^{-\frac{1}{2\alpha^2}\|\mathbf{x}-\mathbf{x}_j\|^2} 1(y_j = y)$$
- The weighted nearest neighbor rule can therefore be interpreted as implementing the MPE rule (= MAP rule) with the above non-parametric estimate of the joint distribution  $p(\mathbf{x}, y)$
- Other kernel fns.:  $\psi(t)$  any “nice” symmetric pdf, nonincreasing in  $|t|$

# Curse of dimensionality

- NN classifiers break down if  $d$  is large
- Suppose training points uniformly distributed in  $d$ -dimensional unit sphere centered at origin
- To “capture” a fraction  $f$  of all training points, the distance  $r$  from the origin that one needs to travel is given by:  $\frac{r^d}{1^d} = f \Rightarrow r = f^{\frac{1}{d}}$
- Say  $d = 10$  (modest).
  - $f = 0.1 \Rightarrow r = 0.8$ , i.e., to capture 10% of samples, need to traverse 80% of range! Samples that are so far from a point (non-local) are typically not good predictors of behavior at that point
  - $f = 0.01 \Rightarrow r = 0.63$ , i.e., to capture only 1% of samples, still need to traverse 63% of range!
  - Seen another way, to keep decision making local, if we fix  $r = 0.1 \Rightarrow$  number of points used for classification  $= nr^d = n10^{-10} \Rightarrow$  will need a HUMONGOUS amount of training data to make a reliable decision

# Curse of dimensionality



- Turns out that in high dimensions, most points are far from each other!
- Also turns out that in high dimensions, most points are almost orthogonal to each other!! ...Weird things happen in high dimensions